

DynamoDB Design for Skier Ride Data Processing

<https://github.com/linmeng128/CS6650-Assignment3>

Database Design Overview

After evaluating several database options (Redis, MySQL/RDS, DynamoDB, MongoDB), we selected **DynamoDB** for the following reasons:

- **Write Throughput:** Superior performance for high-volume write operations (crucial for our message processing)
 - **Scalability:** Ability to handle massive throughput without complex sharding strategies
 - **GSI Support:** Built-in indexing for efficient query access patterns without write locks
 - **Managed Service:** No operational overhead for scaling, replication, or failover
- Integration:** Native AWS integration with our EC2-based consumer architecture

Table Structure

Primary Table: SkierRides

- Partition Key: skierId (Number)
- Sort Key: sortKey (String) – format: "dayId#liftId#timestamp"
- Attributes: resortId, dayId, liftId, time, vertical

Global Secondary Indexes (GSIs)

resort-day-index:

- Partition Key: resortId
- Sort Key: dayId
- Projected Attributes: skierId
- Purpose: Count unique skiers per resort/day

skier-day-index:

- Partition Key: skierId
- Sort Key: dayId
- Projected Attributes: vertical, liftId
- Purpose: Analyze skier activity by day

Deployment Topology

Compute Resources

- **Consumer Application:** Java application running on EC2

- Instance Type: t3.medium
- Role: Consumes messages from RabbitMQ, processes data, writes to DynamoDB
- Concurrency: 512 threads, 250 prefetch count per thread

Database Provisioning

- **DynamoDB Table:** Provisioned capacity mode

- Base Table: 10 RCU, 2000 WCU
- GSIs: 20 RCU, 2000 WCU each
- Region: US-West-2 (Oregon)

Message Queue

- **RabbitMQ:** Deployed on EC2

- Instance Type: t2.micro
- Configuration: Persistent queues, durable messages

Performance Optimization

- **Batch Writing:** Implemented custom batching with 25 items per batch
- **Write Efficiency:** Combined flush interval of 100ms with maximum batch size
- **GSI Scaling:** Matched GSI write capacity to base table (2000 WCU)
- **HTTP Client Tuning:** MaxConcurrency=250, ConnectionTimeout=5s

Client performance metrics:

Phase 1 Results:

=====

Thread count: 32

Total requests: 32000

Successful requests: 32000

Failed requests: 0

Total time: 33.43 seconds

Throughput: 957.37 requests/second

=====

Phase 2 Results:

=====

Thread count: 300

Total requests: 168000

Successful requests: 168000

Failed requests: 0

Total time: 25.44 seconds

Throughput: 6602.74 requests/second

=====

Overall Results:

=====

Phase 1: 32 threads with 1000 requests each

Phase 2: 300 threads with remaining requests

Total Requests: 200000

Successful requests: 200000

Failed requests: 0

Total time: 58.87 seconds

Overall Throughput: 3397.37 requests/second

=====

=== Overall Performance Metrics ===

Sample size: 200000 requests

Mean Response Time: 40.89 ms

Median Response Time: 36 ms

Min Response Time: 12 ms

Max Response Time: 470 ms

90th Percentile Response Time (p90): 60 ms

95th Percentile Response Time (p95): 76 ms

99th Percentile Response Time (p99): 112 ms

Throughput: 3397.37 requests/sec

=== Phase 1 Performance Metrics ===

Sample size: 32000 requests

Mean Response Time: 32.47 ms

Median Response Time: 30 ms

Min Response Time: 12 ms

Max Response Time: 469 ms

90th Percentile Response Time (p90): 38 ms

95th Percentile Response Time (p95): 45 ms

99th Percentile Response Time (p99): 86 ms

=== Phase 2 Performance Metrics ===

Sample size: 168000 requests

Mean Response Time: 42.50 ms

Median Response Time: 37 ms

Min Response Time: 14 ms

Max Response Time: 470 ms

90th Percentile Response Time (p90): 63 ms

95th Percentile Response Time (p95): 79 ms

99th Percentile Response Time (p99): 114 ms

The screenshot shows an IDE with a project named 'client1'. The project structure includes a 'main' directory with a 'java' subdirectory containing 'org.example'. The 'org.example' package contains several classes: 'LiftRideEvent', 'LiftRideGenerator', 'Main', 'PostingThread', 'PostingThreadCSV', 'PostingThreadWithPooling', 'SingleClient', and 'SkierClient'. The 'SkierClient.java' file is open, showing a public class with a main method. The main method calls 'phase1Completed.get()' and 'phase2Completed.get()' to check the completion of two phases. It then calls 'phase1Executor.shutdownNow()' and 'phase2Executor.shutdownNow()' to shutdown the executors. Finally, it calls 'monitorExecutor.shutdownNow()' to shutdown the monitor executor. The 'Run' button is highlighted, and the 'Run' output window shows the following performance results:

```
Phase 1 Results:
=====
Thread count: 32
Total requests: 32000
Successful requests: 32000
Failed requests: 0
Total time: 33.43 seconds
Throughput: 957.37 requests/second
=====
```

The status bar at the bottom indicates the file is 'SkierClient.java' at line 191, column 32, with a UTF-8 encoding and 4 spaces.

```
client1 main pom.xml (Client) SkierClient.java LiftRideGenerator.java PostingThreadWithPooling.java
Project
  .idea
  Client
Run SkierClient
  Phase 2 Results:
  =====
  Thread count: 300
  Total requests: 168000
  Successful requests: 168000
  Failed requests: 0
  Total time: 25.44 seconds
  Throughput: 6602.74 requests/second
  =====
  Overall Results:
  =====
  Phase 1: 32 threads with 1000 requests each
  Phase 2: 300 threads with remaining requests
  Total Requests: 200000
  Successful requests: 200000
  Failed requests: 0
  Total time: 58.87 seconds
  Overall Throughput: 3397.37 requests/second
  =====
  === Overall Performance Metrics ===
  Sample size: 200000 requests
  Mean Response Time: 40.89 ms
  client1 Client > src > main > java > org > example > SkierClient > main 191:32 LF UTF-8 4 spaces
```

```
client1 main pom.xml (Client) SkierClient.java LiftRideGenerator.java PostingThreadWithPooling.java
Project
  .idea
  Client
Run SkierClient
  === Overall Performance Metrics ===
  Sample size: 200000 requests
  Mean Response Time: 40.89 ms
  Median Response Time: 36 ms
  Min Response Time: 12 ms
  Max Response Time: 470 ms
  90th Percentile Response Time (p90): 60 ms
  95th Percentile Response Time (p95): 76 ms
  99th Percentile Response Time (p99): 112 ms
  Throughput: 3397.37 requests/sec
  === Phase 1 Performance Metrics ===
  Sample size: 32000 requests
  Mean Response Time: 32.47 ms
  Median Response Time: 30 ms
  Min Response Time: 12 ms
  Max Response Time: 469 ms
  90th Percentile Response Time (p90): 38 ms
  95th Percentile Response Time (p95): 45 ms
  99th Percentile Response Time (p99): 86 ms
  === Phase 2 Performance Metrics ===
  Sample size: 168000 requests
  Mean Response Time: 42.50 ms
  Median Response Time: 37 ms
  Min Response Time: 14 ms
  Max Response Time: 470 ms
  client1 Client > src > main > java > org > example > SkierClient > main 191:32 LF UTF-8 4 spaces
```

Overview

